

Assessment 3

Waikei Lau (Z5349878)

- [2 points] are given if the submitted R code works with no error.
- [2 points] Read the dataset into R. Check if there are missing values (NA) and, in case there are, remove them.

```
wines <- read.csv("D:\\General\\Bayesian Inference\\winequality-red.csv")
wines <- na.omit(wines)
```

- [2 points] We want to implement a logistic regression, therefore we want a response variable which assume values either 0 or 1. Suppose we consider "good" a wine with quality above 6.5 (included).

```
wines$Good <- ifelse(wines$quality >= 6.5 , 1 , 0)
```

Created new column where wines with score greater than or equal to 6.5 are 1 otherwise 0.

- [4 points] Run a frequentist analysis on the logistic model, using the glm() function. What are the significant coefficients?

```
obj_log <- glm(wines$Good ~ . - quality, family = binomial(link="logit"), data=wines)
summary(obj_log)
```

Using the glm function, implemented a logistic regression model with wine quality as response variable and all other coefficients included in model.

```
Call:
glm(formula = wines$Good ~ . - quality, family = binomial(link = "logit"),
    data = wines)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.9878  -0.4351  -0.2207  -0.1222   2.9869

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    2.428e+02  1.081e+02   2.247  0.024660 *
fixed.acidity    2.750e-01  1.253e-01   2.195  0.028183 *
volatile.acidity -2.581e+00  7.843e-01  -3.291  0.000999 ***
citric.acid      5.678e-01  8.385e-01   0.677  0.498313
residual.sugar   2.395e-01  7.373e-02   3.248  0.001163 **
chlorides       -8.816e+00  3.365e+00  -2.620  0.008788 **
free.sulfur.dioxide 1.082e-02  1.223e-02   0.884  0.376469
total.sulfur.dioxide -1.653e-02  4.894e-03  -3.378  0.000731 ***
density         -2.578e+02  1.104e+02  -2.335  0.019536 *
pH              2.242e-01  9.984e-01   0.225  0.822327
sulphates       3.750e+00  5.416e-01   6.924  4.39e-12 ***
alcohol         7.533e-01  1.316e-01   5.724  1.04e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1269.92  on 1598  degrees of freedom
Residual deviance:  870.86  on 1587  degrees of freedom
AIC: 894.86

Number of Fisher Scoring iterations: 6
```

At $\alpha=0.01$, the significant coefficients are:

Intercept | fixed acidity | volatile acidity | residual sugar | chlorides | total sulfur dioxide | density | sulphates | alcohol

5. [5 points] Estimate the probabilities of having a "success": fix each covariate at its mean level and compute the probabilities for a wine to score "good" varying total.sulfur.dioxide and plot the results.

```
lr_coef <- as.numeric(coefficients(obj_log))
range.tsd <- seq(from=min(wines$total.sulfur.dioxide), to=max(wines$total.sulfur.dioxide), by=1)

b0 <- lr_coef[1]
bx1 <- lr_coef[2]*colMeans(wines)[1]
bx2 <- lr_coef[3]*colMeans(wines)[2]
bx3 <- lr_coef[4]*colMeans(wines)[3]
bx4 <- lr_coef[5]*colMeans(wines)[4]
bx5 <- lr_coef[6]*colMeans(wines)[5]
bx6 <- lr_coef[7]*colMeans(wines)[6]
bx7 <- lr_coef[8]*range.tsd
bx8 <- lr_coef[9]*colMeans(wines)[8]
bx9 <- lr_coef[10]*colMeans(wines)[9]
bx10 <- lr_coef[11]*colMeans(wines)[10]
bx11 <- lr_coef[12]*colMeans(wines)[11]


t.s.d.Good <- b0+bx1+bx2+bx3+bx4+bx5+bx6+bx7+bx8+bx9+bx10+bx11

Prob.Good.wine <- exp(t.s.d.Good)/(1+exp(t.s.d.Good))

plot(range.tsd,Prob.Good.wine,ylim=c(0,1), type="l", lwd=2, col="blue",
      xlab="Total Sulfur-dioxide", ylab="P(Good Wine)", main="Probability of Good wine")

abline(h=.5, lty=2)
```

Fixing all covariates at mean levels, and multiplying by the corresponding MLE from logistic regression, we see that the probability of a good wine increases to ~10% as Total Sulfur dioxide decreases.



6. [15 points] Perform a Bayesian analysis of the logistic model for the dataset, i.e. approximate the posterior distributions of the regression coefficients, following these steps:

- Write an R function for the log posterior distribution.
- Fix the number of simulation at 10^4 .
- Choose 4 different initialisations for the coefficients.
- For each initialisation, run a Metropolis–Hastings algorithm.
- Plot the chains for each coefficients (the 4 chains on the same plot) and comment.

```

# Log-posterior distribution
lpost.LR <- function(beta,x,y){
  eta <- as.numeric(x %*% beta)
  logp <- eta - log(1+exp(eta))
  logq <- log(1-exp(logp))
  logl <- sum(logp[y==1]) + sum(logq[y==0])
  lprior <- sum(dnorm(beta,0,10,log=T))
  return(logl + lprior)
}

S <- 10^4

X = as.matrix(cbind(rep(1,nrow(wines)), wines[,1:11]))
colnames(X)[1] <- "Intercept"
y = wines[, "Good"]

k <- ncol(X)
acc <- 0
Omega_prop <- solve(t(X) %*% X)

library(mvtnorm)

beta_mat1 <- beta_mat2 <- beta_mat3 <- beta_mat4 <- matrix(NA, nrow=S, ncol= ncol(X))
colnames(beta_mat1) <- colnames(beta_mat2) <- colnames(beta_mat3) <- colnames(beta_mat4) <- colnames(X)

beta_mat1[1,] <- lr_coef
beta_mat2[1,] <- c(10,0.5,-3, 1.2,0.3,-10,0.05,-0.018,-10,0.3,4,0.9)
beta_mat3[1,] <- c(-10,-0.1,-5,-1,-0.35,-12,0.05,-0.01,10,0,0,0)
beta_mat4[1,] <- c(rep(0,12))

tuning = 0.85
acc=0

```

```

for(i in 2:S){

  # 1. Propose a new set of values
  beta_star1 <- rmvnorm(1,beta_mat1[i-1,],tuning*Omega_prop)
  beta_star2 <- rmvnorm(1,beta_mat2[i-1,],tuning*Omega_prop)
  beta_star3 <- rmvnorm(1,beta_mat3[i-1,],tuning*Omega_prop)
  beta_star4 <- rmvnorm(1,beta_mat4[i-1,],tuning*Omega_prop)

  # 2. Compute the posterior density on the proposed value and on the old value
  newpost1=lpost.LR(t(beta_star1),X,y)
  newpost2=lpost.LR(t(beta_star2),X,y)
  newpost3=lpost.LR(t(beta_star3),X,y)
  newpost4=lpost.LR(t(beta_star4),X,y)

  oldpost1=lpost.LR(matrix(beta_mat1[i-1,],ncol=1),X,y)
  oldpost2=lpost.LR(matrix(beta_mat2[i-1,],ncol=1),X,y)
  oldpost3=lpost.LR(matrix(beta_mat3[i-1,],ncol=1),X,y)
  oldpost4=lpost.LR(matrix(beta_mat4[i-1,],ncol=1),X,y)

  # 3. Acceptance step
  if(runif(1,0,1)>exp(newpost1-oldpost1)){
    beta_mat1[i,]=beta_mat1[i-1,]
  } else {
    beta_mat1[i,]=beta_star1
    acc=acc+1
  }
  if(runif(1,0,1)>exp(newpost2-oldpost2)){
    beta_mat2[i,]=beta_mat2[i-1,]
  } else {
    beta_mat2[i,]=beta_star2
    acc=acc+1
  }
  if(runif(1,0,1)>exp(newpost3-oldpost3)){
    beta_mat3[i,]=beta_mat3[i-1,]
  } else {
    beta_mat3[i,]=beta_star3
    acc=acc+1
  }
  if(runif(1,0,1)>exp(newpost4-oldpost4)){
    beta_mat4[i,]=beta_mat4[i-1,]
  } else {
    beta_mat4[i,]=beta_star4
    acc=acc+1
  }

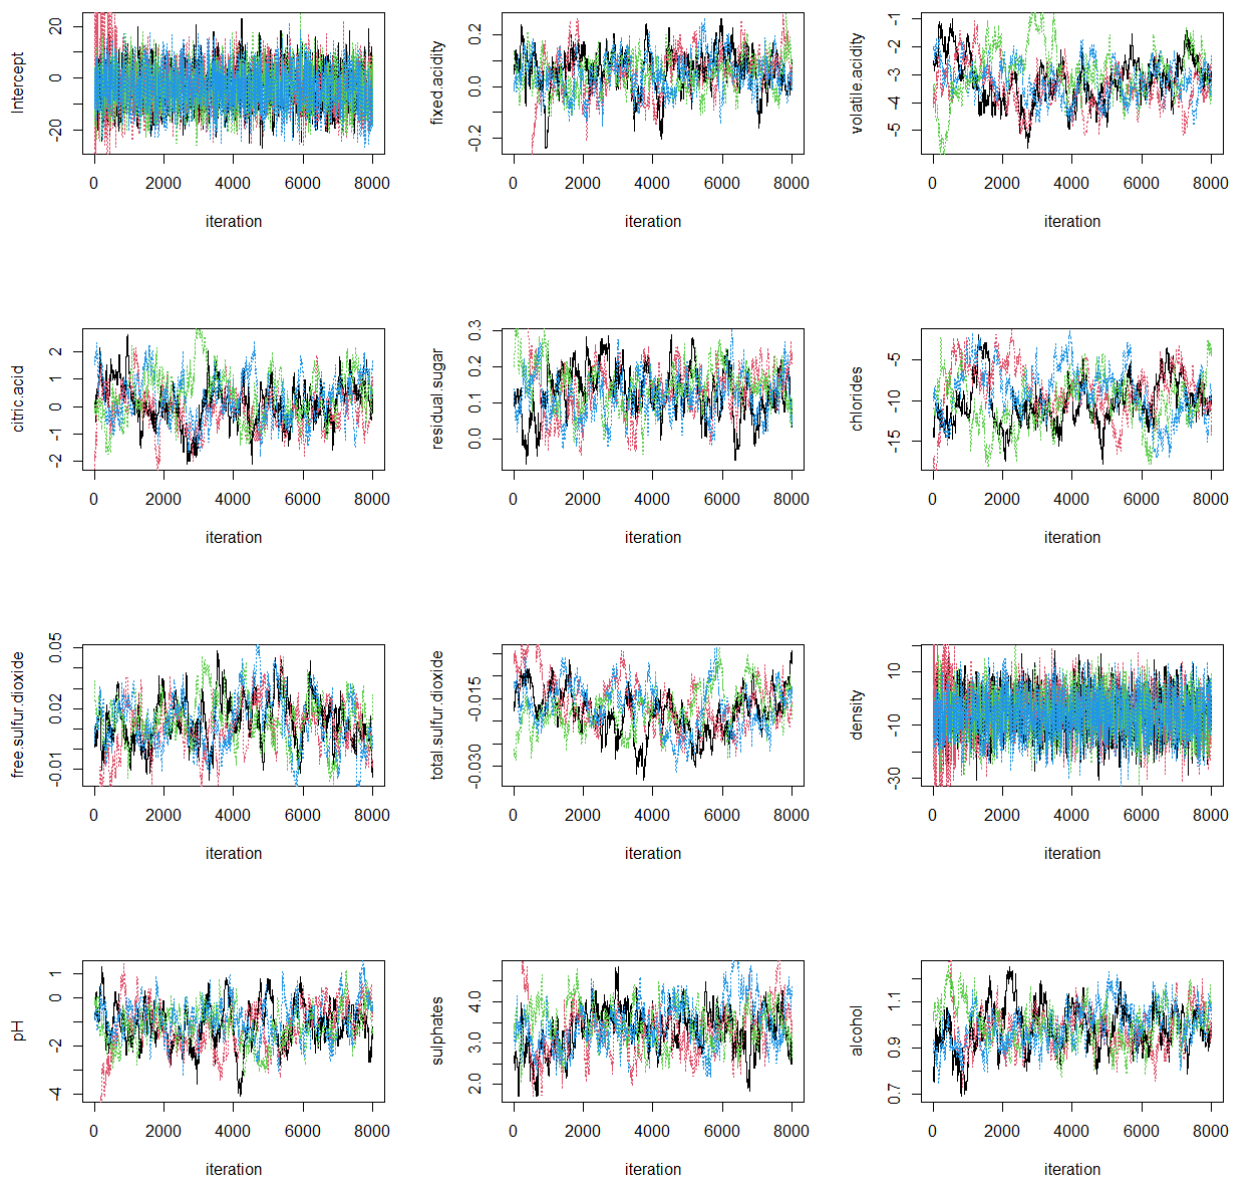
  # 4. Print the stage of the chain
  if(i%%1000==0){print(c(i,acc/(i*4)))}
}

```

```

burn.in= 2000
par(mfrow=c(4,3))
for(i in 1:12){
  plot(beta_mat1[,i],type="l", col=1, ylab=colnames(beta_mat1)[i], xlab = "iteration")
  lines(beta_mat2[,i],type="l", col=2, lty=3)
  lines(beta_mat3[,i],type="l", col=3, lty=3)
  lines(beta_mat4[,i],type="l", col=4, lty=3)
}
par(mfrow=c(4,3))
for(i in 1:12){
  plot(beta_mat1[burn.in:S,i],type="l", col=1, ylab=colnames(beta_mat1)[i], xlab = "iteration")
  lines(beta_mat2[burn.in:S,i],type="l", col=2, lty=3)
  lines(beta_mat3[burn.in:S,i],type="l", col=3, lty=3)
  lines(beta_mat4[burn.in:S,i],type="l", col=4, lty=3)
}

```



The log posterior distribution is:

$$\log L(\beta; \mathbf{y}, \mathbf{x}) = \log \left[\prod_{i=1}^n \left(\frac{\exp(\mathbf{x}_i \beta)}{1 + \exp(\mathbf{x}_i \beta)} \right)^{y_i} \left(1 - \frac{\exp(\mathbf{x}_i \beta)}{1 + \exp(\mathbf{x}_i \beta)} \right)^{1-y_i} \right]$$

Four different initialisations were chosen:

- Using MLE from logistic regression model
- Two randomly selected starting points based on deviations from MLE and standard deviations.
- All coefficients starting from zero.

A tuning parameter of 0.85 was selected for an acceptance rate close to 23%.

A burn-in was selected at 2,000 out of 10,000 – where in most coefficients appear to have achieved convergence.

Observations on the intercept and the density coefficient that deviated from the MLE of 242.8 and -257.8 respectively. Convergence was observed around -3.29 and -5.1 respectively with no autocorrelation detected.

7. [5 points] Approximate the posterior predictive distribution of an unobserved variable characterised by:

- fixed acidity: 7.5 | volatile acidity: 0.6 | citric acid: 0.0 | residual sugar: 1.70 | chlorides: 0.085 | free sulfur dioxide: 5 | total sulfur dioxide: 45 | density: 0.9965 | pH: 3.40 | sulphates: 0.63 | alcohol: 12

Plot the approximate posterior predictive distribution.


```
y_new <- c(1)
x_new <- c(1,7.5,0.6,0,1.7,0.085,5,45,0.9965,3.4,0.63,12)
par(mfrow=c(1,1))

for(i in burn.in:S){

  # 5. Prediction
  p_new <- exp(sum(beta_mat1[i,] * x_new) ) / (1 + exp(sum(beta_mat1[i,] * x_new) ))
  y_new[i] <- rbinom(1,1,prob=p_new)

}

hist(y_new, main="Post-predictive distribution",
     breaks = 2, col="green")
```



Implemented predictive X values in combination with the beta matrix to derive probability of success. This was incorporated into Bernoulli (binomial) samples on the response variable y.

Total accepted: 831

Total rejected: 7171

Acceptance rate: 10.38%

8. [5 points] Use the `metrop()` function available in the `mcmc` package to perform the same analysis on the posterior distribution you have approximated for Question 6. Choose again 10^4 simulations and compare the results with the results obtained with your code. (Here a visual comparison of the chains is enough to get full mark).

```
# Log-posterior distribution
lpost.metr <- function(x,y)function(beta){
  eta <- as.numeric(x %%% beta)
  logp <- eta - log(1+exp(eta))
  logq <- log(1-exp(logp))
  logl <- sum(logp[y==1]) + sum(logq[y==0])
  lprior <- sum(dnorm(beta,0,10,log=T))
  return(logl + lprior)
}

lpost <- lpost.metr(X, y)

burn.in= 2000
tuning = 0.25

out1 <- metrop(lpost, beta_mat1[,1], S, scale=tuning*Omega_prop)
out2 <- metrop(lpost, beta_mat2[,1], S, scale=tuning*Omega_prop)
out3 <- metrop(lpost, beta_mat3[,1], S, scale=tuning*Omega_prop)
out4 <- metrop(lpost, beta_mat4[,1], S, scale=tuning*Omega_prop)

out1$accept
out2$accept
out3$accept
out4$accept

par(mfrow=c(4,3))

axis <- rbind(c(-25,25),c(-0.5,0.8),c(-6,1),c(-1.8,2),c(-0.5,0.5),c(-15,1),
              c(-0.02,0.07),c(-0.025,0.005),c(-25,25),c(-5,5),c(-1,5),c(-0.2,1.2))

for(i in 1:12){
  plot(out1$batch[burn.in:S,i],type="l", col=1, ylab=colnames(beta_mat1)[i], xlab = "iteration", ylim=axis[i,])
  lines(out2$batch[burn.in:S,i],type="l", col=2, lty=3)
  lines(out3$batch[burn.in:S,i],type="l", col=3, lty=3)
  lines(out4$batch[burn.in:S,i],type="l", col=4, lty=3)
}
```

Using the `metrop()` function, both the intercept and density coefficients appear to have achieved the stationary distribution where all four initialisations converged. However, other coefficients did not achieve convergence within 10,000 iterations using the `metrop()` function.

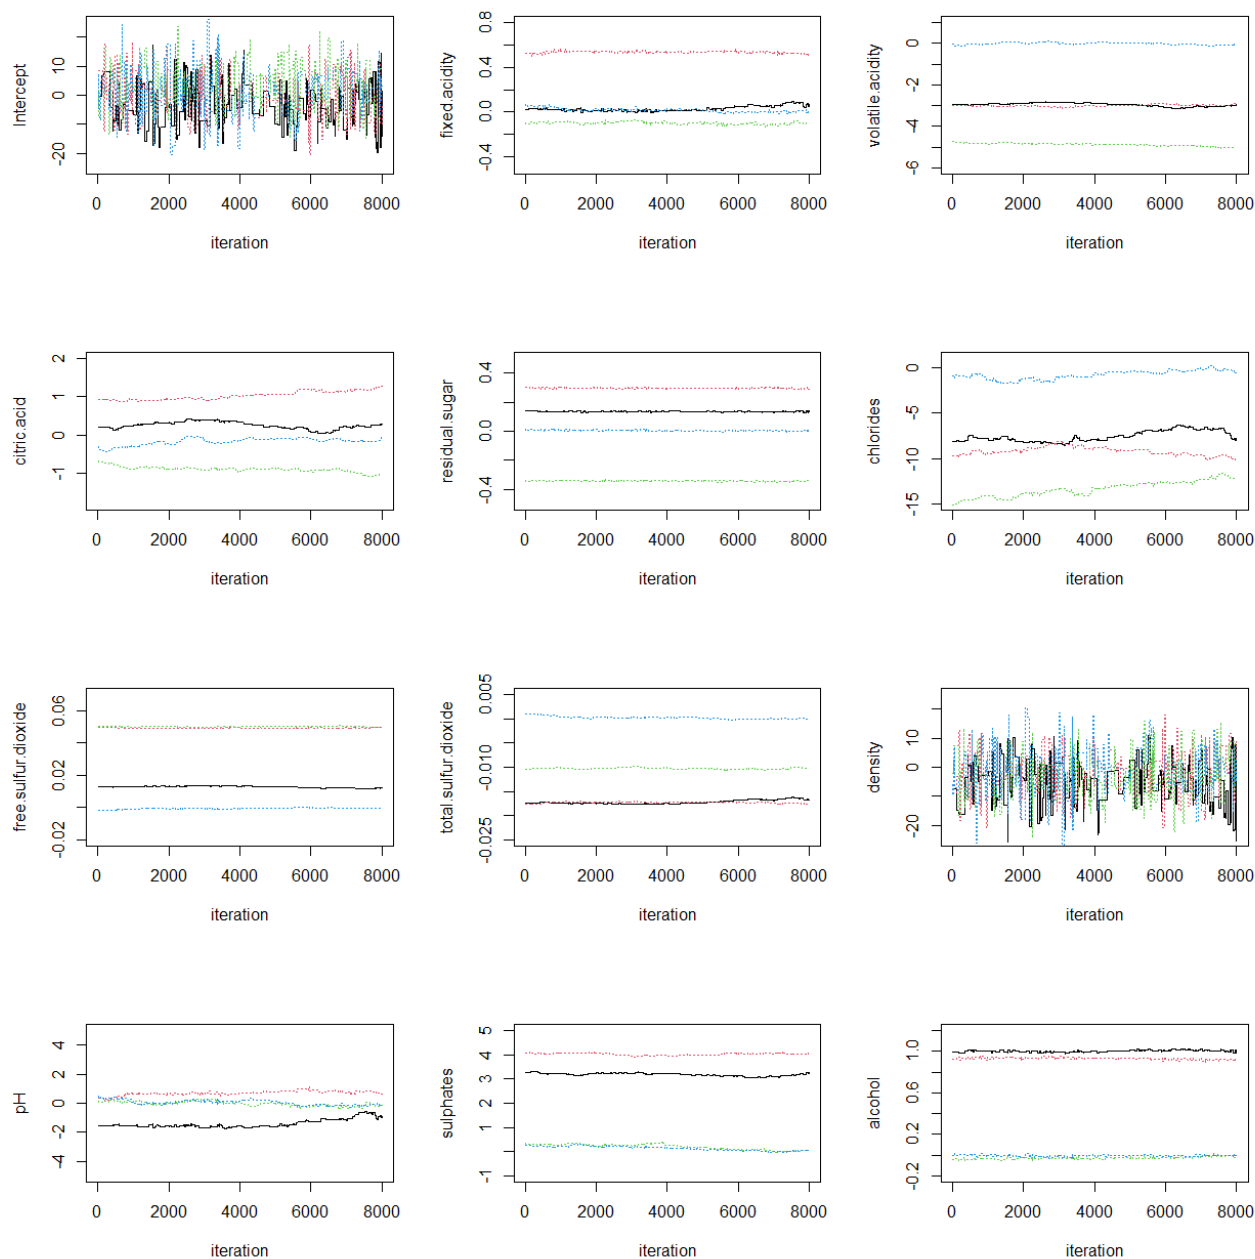
As can be seen from the plots after a burn in of 2,000 observations was applied, the plots of various initialisations did not converge. This was after using the four initialisations from earlier:

- Using MLE from logistic regression model
- Two randomly selected starting points based on deviations from MLE and standard deviations.
- All coefficients starting from zero.

A smaller tuning parameter of 0.25 was required to achieve an acceptance rate close to 23%, which was gradually adjusted to 0.9 to test for convergence. The plots are provided below showing separation between coefficients at both 25% and 90% tuning.

Additionally there appears to be autocorrelation as the coefficients follow a relatively flat trend for consecutive iterations indicating multiple rejections.

Tuning at 25%



Tuning at 90%

